

Multiprocessing in ArcPy

Esri Dev Summit 2016

Bryan Chastain

<https://github.com/bchastain/devsummit2016>

1

CGI

The Problem

- You have some complicated geoprocessing task that needs to be repeated over a large number of inputs
 - Exporting 1000 maps for a project
 - Performing the same hydraulic raster calculations on 500 different DEMs
 - Performing sensitivity analysis on a spatial simulation
 - Running MCMC methods on spatial Bayesian models

2

CGI


Simple example: export 100 maps

```
import arcpy

mxd = arcpy.mapping.MapDocument('./testmap.mxd')
out = './images/out'
for i in (0,100):
    arcpy.mapping.ExportToJPEG(mxd, '%s%s.jpg' % (out, i), '',
                               1056, 816, 96, '',
                               '8-BIT_GRAYSCALE', 100)
```

CGI

Despite mad Python skills, you still see this:




Script runtime:
375.1s

CGI

I know what to do!

- Being a smart coder, familiar with similar handling in C++, Java, etc., you think: **Multithreading!**
- Gives ability to use the multiple cores in your machine to concurrently execute multiple tasks in parallel




5

CGI

Hey, cool, a Threading module in Python!

- Easy to use, simply extend the `threading.Thread` class and override its `run()` function to set up the code to be parallelized
- The `Queue` module is also a handy counterpart for delivering work to your new `Thread` class and flagging it as done



6

CGI

Exporting again, but with multithreading

```
from Queue import Queue
from threading import Thread
import arcpy

class DownloadWorker(Thread):
    def __init__(self, queue):
        Thread.__init__(self)
        self.queue = queue
    def run(self):
        while True:
            # Get the work from the queue and expand the tuple
            filename = self.queue.get()
            mxd = arcpy.mapping.MapDocument(r'C:\CGI\presentations\testmap.mxd')
            arcpy.mapping.ExportToJPEG(mxd, '',
                                      1056, 816, 96, '', '8-BIT_GRAYSCALE', 100)
            self.queue.task_done()
```




Exporting again, but with multithreading

```
if __name__ == '__main__':
    out = './images/out'
    filenames = []
    for i in range(0,100):
        filenames.append('%s%s.jpg' % (out, i))
    # Create a queue to communicate with the worker threads
    queue = Queue()
    # Create 8 worker threads
    for x in range(8):
        worker = DownloadWorker(queue)
        # Setting daemon to True will let the main thread exit
        worker.daemon = True
        worker.start()
    # Put the tasks into the queue as a tuple
    for file in filenames:
        queue.put(file)
    # Causes the main thread to wait for the queue to finish
    queue.join()
```



Python multithreading in action



Script runtime:
713.3s

CGI

Hey wait, what?!

- CPython's Global Interpreter Lock (GIL)
 - Protects data from being accessed by multiple threads (e.g. current thread state and heap allocated objects for garbage collection)
 - Has the effect of limiting Python to executing one **single thread** at a time
 - For more, read Jeff Knupp's "Python's Hardest Problem"

<https://jeffknupp.com/blog/2012/03/31/python-hardest-problem/>

Everything I Know About Python...

The personal blog of author, speaker, tutor, and professional software engineer Jeff Knupp



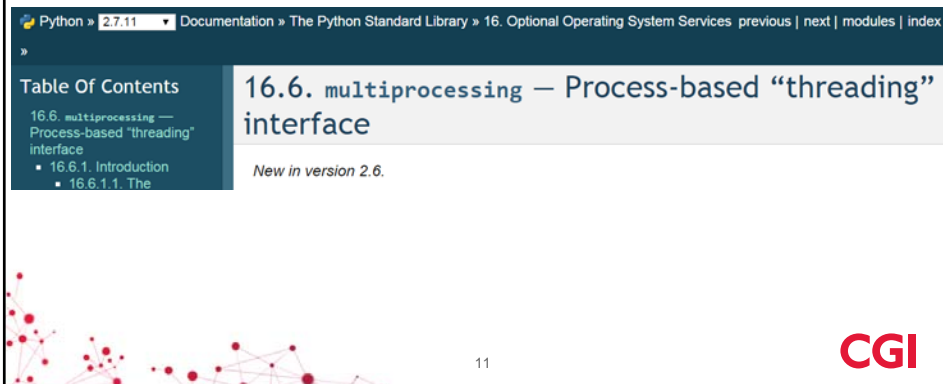
Python's Hardest Problem

For more than a decade, no single issue has caused more frustration or curiosity for Python novices and experts alike than the Global Interpreter Lock.

CGI

Enter the multiprocessing module

- GIL limits us to 1 *thread* at a time, but no limit on *processes*
- Starting in Python 2.6, the multiprocessing module lets you write parallelized code that bypass GIL issues



Python » 2.7.11 » Documentation » The Python Standard Library » 16. Optional Operating System Services previous | next | modules | index

»

Table Of Contents

16.6. multiprocessing — Process-based “threading” interface

- 16.6.1. Introduction
- 16.6.1.1. The

16.6. multiprocessing — Process-based “threading” interface


New in version 2.6.

11

CGI

Usage, in simplest case

- `multiprocessing.Pool & map()`
 - For simple cases, where no synchronization is needed between processes and the tasks are truly independent of one another
- The Pool class represents a pool of worker processes
 - By default is equal to number of CPUs available, but can be set to any number
- `pool.map(func, iterable)`
 - Parallel equivalent of built-in Python `map()` function
 - Chops the iterable into a number of chunks which it submits to the process pool as separate tasks



12

CGI


Example

```
import multiprocessing as mp
import time
import arcpy
# Function to map
def exportmap(filename):
    mxd = arcpy.mapping.MapDocument('./testmap.mxd')
    arcpy.mapping.ExportToJPEG(mxd, filename, "", 1056, 816, 96,
                               "", "8-BIT_GRAYSCALE", 100)

if __name__ == '__main__':
    # Optional, default value anyways
    NUM_PROCESSES = mp.cpu_count()
    pool = mp.Pool(NUM_PROCESSES)
    out = './images/out'
    # Create list of filenames as our iterable
    filenames = ['%s%s.jpg' % (out, i) for i in range(100)]
    pool.map(exportmap, filenames)
```

CGI

That's more like it...



Script runtime:
127.8s

CGI

More multiprocessing

- `pool.map_async(func, iterable)`
 - Same as `map`, but results are returned asynchronously (does not block)
 - Returns `pool.AsyncResult`
- `pool.apply(func, args)`
 - Similar to `pool.map()`, but only spawns a single worker (would need to call multiple times to generate multiple processes)
- `pool.apply_async(func, args)`
 - Same as `apply()`, but results are returned asynchronously (does not block)
 - Returns `pool.AsyncResult`

15

CGI

Advanced multiprocessing

- `mp.Process` class
 - Instead of working w/ Pool of workers, direct control over each process
- Synchronization between processes
 - `mp.Lock.acquire()` & `mp.Lock.release()`
 - Acquire/release locks on std-out & files to prevent jumble of output
- Communication between processes
 - `mp.Pipe`
 - For sending/receiving pickleable objects between processes
- Manage/control processes on different computers across a network
 - `mp.Manager`

16

CGI

Other considerations

- `partial()`
 - `pool.map()` & `map_async()` only accept one parameter for their “func”, the iterable
 - If you want static parameters to be passed in addition to the parallelized iterable, need to use `partial()`

```
def myfunc(a, b):  
    print '%s: %i' % (a, b)  
if __name__ == '__main__':  
    the_func = partial(myfunc, 'static text')  
    pool.map(the_func, range(100))
```

17

CGI

Other considerations

- Try to use “in_memory” workspace for temporary data.
 - Can improve performance over writing data to disk.
 - However, size of data may prevent this.
 - Deleting in-memory dataset when finished can prevent memory errors

18

CGI

Limitations

- Processes have considerably higher initialization overhead than threads
 - May not make sense to use in cases where # of tasks to be parallelized is small or work within each task is quick
 - Consider the number of processes in a pool
 - Note: multiprocessing our export maps script across 24 CPUs did not come anywhere close to a 24x speed-up
- Tasks to parallelize should be independent as possible and non-serial
- Need to consider GDB schema locks
- Possible to run w/i ArcGIS Desktop, but performance may suffer
 - Also, need to un-check "run script in process"
- Debugging can be difficult

19

CGI

Resources


- Reading on the GIL
 - <https://jeffknupp.com/blog/2012/03/31/pythons-hardest-problem/>
 - <https://wiki.python.org/moin/GlobalInterpreterLock>
- Esri multiprocessing samples
 - <https://blogs.esri.com/esri/arcgis/2012/09/26/distributed-processing-with-arcgis-part-1/>
- Distributed/cluster computing with IPython
 - <http://ipython.org/ipython-doc/stable/parallel/index.html>

20

CGI

Questions?

<https://github.com/bchastain/devsummit2016>



21

CGI